

Detecting child unsafe videos using transcripts

Shritishma Reddy & Pratik Kamble

April, 16th 2019

1 Abstract

In this project, we propose an end-to-end pipeline for the detection of child unsafe videos using video transcripts. Given the spread of digital media and increase in media consumption by kids (ages 3-12) over the past decade, there is a great need for the filtering of unsafe content found within reach of said age group. While traditional parental systems do exist, this would be an online scalable system, that would only benefit from the increasing video consumption.

2 Introduction

This pipeline has three phases:

- Data acquisition: Collection of a holistic video-transcript dataset representative of children viewing habits and trends.
- Unsafe pattern detection: Analyze speech patterns and clustering hotspot zones to form a timeline representation.
- Video Classification: Classification of videos into unsafe types, if unsafe, viz. religious, abusive, drug use, etc.

2.1 Data Acquisition

The data acquisition process depends a lot on the type of genre distribution. A study [3] has found that children consume just over three hours of media, as of 2015; the numbers having since then gone up further.

The time spent with on screen media dramatically increases from the toddler to preschool to school-age years. Children under two have a screen time average of 53 minutes per day. This increases to almost two and a half hours per day among two to four year old and almost three hours for kids in the five to eight year old range. By age eight, 96% of children have watched TV, 90% have used a computer, 81% have played console video games, and 60% have played games or used apps on a portable device. Thus by order of data-source, we should

have a dataset representative mainly of TV show content, followed by YouTube and lastly, game video data.

Luckily, all video footage from various sources can be obtained from YouTube, along with transcript data. TV show snippets, actual online vlog and other consumable media, as well as video game footage can be found on the streaming platform.

2.2 Unsafe Pattern Detection

Once we have the annotated data, we have to run through it and classify according to unsafe-ness. One criterion that is of importance is categories that may influence minds may be quite different from ones for adolescents or adults.

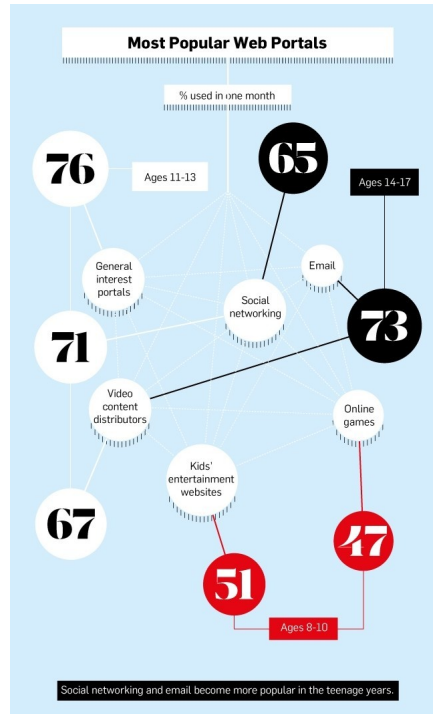


Figure 1: Major content sources and distribution

As seen from the figure , the main categories of media streaming are entertainment websites, online games, email, social networking and general interest portals. Breaking down the type of content found on these sites, the main type of genres are as follows:

- Religion-based hate speech: This occurs in chatrooms and online forums, both high target kid influence
- Color-based hate speech: Depending on race and ethnicity, such sort of speech may be hidden in meaning and harder to detect.
- Foul Language: Kids usually pick up on foul language and the age of such exposure is lower this decade than it ever was.
- Other offensive terminology: Dealing with kids in the lower age bracket, systems need not worry too much about false positives, because of the ease of influence of media on early youth.

2.3 Video Classification

Once the data collection for the transcripts is done, a simple classical model is used in cascade, to first analyze and mark scenes and zones in the time frame with dangerous content.

This meta information is then passed to a zonal mapping pipeline, that, depending on how many zones of what clusters there are, estimates a score for each category mentioned prior. Finally, depending on the thresholding for each genre distribution across ages, a weighted average will judge the videos overall scores.

3 Annotation and labeling

The first step in the pipeline is the data collection and annotation. It is important to keep in mind the divisions and genre distribution amidst children and have a dataset representative of that. We are collecting 5 videos in each positive category, approximately of 5 minutes each, unless the genre specifies (eg, vlogs are usually longer). The videos should be at such an age genre balance that younger children usually are pressured into by peers, and boundary videos are easy gateways to explicit content.

This gives us a total of 4 categories of positive videos, for each age group bracket within kids, targetting most common videos in the given set.

$$N = n_{videos} \times n_{categories} \times n_{agegroups} \times m \times \delta \times \alpha$$

Here n_{videos} is 5 per category, $n_{agegroups}$ is 3, m being the average length per video of selection. Now the average rate of speech is δ . This may vary across various demographics, but we choose the global average, 110 words per minute. Let's assume we are targetting positive samples to seed the dataset, there is a 20% conversion rate, α . This means that 1 minute in a positive truth, 5 minute video contains positive samples of the content we're trying to flag. This is a fair assumption, and goes in line with average statistics [2].

3.1 Raw Transcript Collection

After aggregating the various sources of media and content, our dataset consisted of 60 video subtitle files (*.srt). These files were formatted using the srt standard and had to be converted into a ‘DataFrame’ for annotation. After cleaning invalid records and repeat rows, we were left with around 7000 observations.

3.2 SRT Parsing and Dataset Aggregation

Once we had a collection of formatted transcripts, they had to be aggregated and labeled. Each record of the dataset is of the form:

id	genre	idx	start	end	text	
0	3tgZ...	0	1	0.00	3.71	mr. mine you say babe ...
1	3tgZ...	0	3	3.72	22.79	mr. mine babe what’s d...
2	3tgZ...	0	4	22.79	22.80	of a dynamic punch mr. line
3	3tgZ...	0	5	22.80	28.40	of a dynamic punch you’ll...
4	3tgZ...	0	6	28.40	28.41	you’ll get stronger your

Every sample has a start and end marker, the genre for the video, and other meta information (as seen above). The timestamps and genre will help correlate the overall accuracy of the system with certain trends and patterns across genres. Additionally, some inference can be drawn as to the average "unsafe" duration of a video and the subsequent variation with genres.

Now, this annotated data is passed onto the model. The preliminary embedding and tagging of every sentence of a video is handled by the same. The results from said model are then passed onto an aggregator; it takes into account sentence based transcript data, augments it with additional meta information about the video and then makes a prediction regarding the safety of the video in terms of child consumption.

4 Model - Sentence labeling

The model choice we ultimately went with was a bi-LSTM model with a simple embedding layer (from Keras). We used sequence model architectures i.e., BLSTM with character grams as inputs for training. BLSTM and LSTM models usually consists of three sequential layers:

- Embedding Layer
- Bidirectional LSTM
- Fully Connected Layer

To compare the benefits of this model against other classical/hybrid approaches, we pre trained all our models on the Davidson hate speech dataset [?

]. Across the board, the accuracy of the chosen model was better than all other models considered (*a comparative analysis has been provided in the **results** section.*) A small improvement to the pipeline could be considered at this stage, viz. an ensemble approach with a mix of this bi-lstm model and a more classical approach.

4.1 Structure

- To deal with complex structured conversation data, we plan to use word level trigrams to represent each conversation.
- Sequence of trigrams sent through an embedding layer which takes one-hot vector of each character trigram in the sentence and learns the lower dimension representation for it.

4.2 Functionality & Reasoning

- Bi-Directional LSTMs have the ability to capture much richer sequential patterns from both directions of a sequence.
- Each character trigram from the previous embedding layer is fed as a sequence of trigrams, to the BLSTM layer.
- The output of the BLSTM layer is given as input to a Fully Connected (FC) layer which models the interactions between these features.
- The final softmax node in the FC layer outputs the probability of the sentence belonging to the inappropriate class.

4.3 Initial testing & results

The initial training and validation was done on the Davidson dataset. This model gave excellent results, resulting in over 95% training accuracy and over 80% test accuracy on the baseline dataset. Results on the test dataset and comparison across all models is presented in the *results* section.

5 Zonal clustering and overall classification

All the transcript data has been labeled as follows:

- $\begin{bmatrix} 1 & 0 & 0 \end{bmatrix}$ Abusive
- $\begin{bmatrix} 0 & 1 & 0 \end{bmatrix}$ Hate
- $\begin{bmatrix} 1 & 1 & 0 \end{bmatrix}$ Abusive + Hate
- $\begin{bmatrix} 0 & 0 & 1 \end{bmatrix}$ Clean

```

Epoch 21/30
19783/19783 [=====] - 86s 4ms/step - loss: -0.8691 - acc: 0.9391 - val_loss: 0.1019 - val_ac
c: 0.8428
Epoch 22/30
19783/19783 [=====] - 80s 4ms/step - loss: -0.8713 - acc: 0.9400 - val_loss: 0.1204 - val_ac
c: 0.8408
Epoch 23/30
19783/19783 [=====] - 81s 4ms/step - loss: -0.8731 - acc: 0.9404 - val_loss: 0.2340 - val_ac
c: 0.8402
Epoch 24/30
19783/19783 [=====] - 74s 4ms/step - loss: -0.8730 - acc: 0.9403 - val_loss: 0.2093 - val_ac
c: 0.8390
Epoch 25/30
19783/19783 [=====] - 78s 4ms/step - loss: -0.8710 - acc: 0.9394 - val_loss: 0.2144 - val_ac
c: 0.8392
Epoch 26/30
19783/19783 [=====] - 76s 4ms/step - loss: -0.8726 - acc: 0.9400 - val_loss: 0.1851 - val_ac
c: 0.8390
Epoch 27/30
19783/19783 [=====] - 85s 4ms/step - loss: -0.8750 - acc: 0.9404 - val_loss: 0.1948 - val_ac
c: 0.8386
Epoch 28/30
19783/19783 [=====] - 78s 4ms/step - loss: -0.8737 - acc: 0.9403 - val_loss: 0.2013 - val_ac
c: 0.8402
Epoch 29/30
19783/19783 [=====] - 79s 4ms/step - loss: -0.8745 - acc: 0.9403 - val_loss: 0.2292 - val_ac
c: 0.8344
Epoch 30/30
19783/19783 [=====] - 91s 5ms/step - loss: -0.8741 - acc: 0.9400 - val_loss: 0.2117 - val_ac
c: 0.8352

```

Figure 2: bi-LSTM training

This data is then mapped into a timeline heat map to indicate stronger zones of the video that may be unsafe. As part of the further work, this is the step wherein more meta information, age group, region, etc would be added so as to provide a better guess as part of the second phase (*something like an additional FC layer, with the transcript heatmap as a feature subset*).

Now, as part of the first post processing step, we take the labeled video timeline and calculate the percentage of the video that's "unsafe". This is then cross-referenced with the genre based averages and an appropriate genre is determined. Once this part has been done, each genre has a statical threshold for it's unsafeness (again, precalculated). If the video then falls beyond such a threshold then, it is classified as unsafe. Additionally, it is run across other averages as well, and a score is given depending on the sensitivity of the system (*higher sensitivity would lead to more false positives*).

These are then the final labels that are taken into account. As part of the results section, we have compared the end-to-end accuracy with ground truth obtained from manual viewing and judgment.

6 Results

6.1 Data preparation & model training

6.2 Final Choice & Testing

6.3 Pipeline results

7 Conclusion & further work

Given the two step process of "classification", some interesting results can be drawn. From the dataset that we collected, we noticed:

7.1 Further Work

Thus, with the classification methods of the bi-lstm model with a combination of the classical method should give us the best performance. Additionally, information from the classical method can be used in the later clustering and analysis step. A later step is to provide a finer analysis of the type of content and possibly give an appropriate age rating, though this might require some additional work.

Another improvement on the to-do list is to supplement the timeline information with more features, say region and appropriate age group. This way, an online learning model would be able to learn trends in consumption behaviour and auto-adjust to regional trends and filtering.

You can find the complete codebase, along with some helpful scripts on GitHub. Feel free to contribute and request features!

References

- [1] Bi-lstm based tweet hate speech detection. <https://github.com/sebastiandziadzio/hate-tweet/>.
- [2] Biteable. Video statistics 2018. <https://biteable.com/blog/tips/video-marketing-statistics/>, 2018.
- [3] Dr. Brent Conrad. Viewing statistics by time. <http://www.techaddiction.ca/media-statistics.html>, 2015.